

# Kavya Sridhar

San Jose, CA | 858-372-3468 | [kavyasridhar2001@gmail.com](mailto:kavyasridhar2001@gmail.com) | [linkedin.com/in/kavyasridhar/](https://linkedin.com/in/kavyasridhar/) | [github.com/kavyasridhar1501](https://github.com/kavyasridhar1501)

## SUMMARY

Software engineer with production experience owning code from design through deployment and on-call operations, having shipped scalable microservices processing 50M+ daily records at a top-5 European bank. Proficient in Java, Python, and TypeScript with hands-on AWS, SQL/NoSQL, CI/CD, and GenAI-assisted development. AWS Certified AI Practitioner seeking SDE roles building fault-tolerant distributed systems at scale.

## EDUCATION

|                                                                                                   |                                       |
|---------------------------------------------------------------------------------------------------|---------------------------------------|
| <b>University of California, San Diego</b><br><i>Masters in Data Science (GPA 3.92)</i>           | San Diego, CA<br>Sep. 2024 – Jun 2026 |
| <b>College of Engineering Guindy</b><br><i>B.E in Computer Science and Engineering (GPA 3.98)</i> | Chennai, IN<br>Aug. 2019 – Apr 2023   |

## EXPERIENCE

|                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                              |                                             |
|--------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|---------------------------------------------|
| <b>Software Engineer</b><br><i>Dreamline AI</i> <ul style="list-style-type: none"><li>Cut bid lifetime from 4+ months to 48 hours via configurable TTL and PII masking in Supabase</li><li>Blocked 100% of premature project closures via an owner-confirmed, milestone-gated state machine and PostgreSQL schema migration</li></ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                        | Aug 2026 – Present<br>San Jose, CA          |
| <b>Graduate Instructional Assistant</b><br><i>UC San Diego</i> <ul style="list-style-type: none"><li>Decreased manual grading by 65% for 200+ students using a Python autograding pipeline on a Spark cluster</li><li>Recovered 3 hours per assignment by diagnosing failing distributed Spark jobs across scalable-analytics coursework</li><li>Cut regrade requests 30% across 50+ students using reusable R scripts for statistical modeling</li></ul>                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                                    | Sep–Dec 2025, Mar–Jun 2026<br>San Diego, CA |
| <b>Software Engineer</b><br><i>Société Générale</i> <ul style="list-style-type: none"><li>Eliminated 100% of manual alert triage via Foundry dashboards classifying 50M+ daily records into severity tiers</li><li>Accelerated daily AML batch ETL 40% via Spark transformations and Parquet partitioning on AWS S3 &amp; Hadoop</li><li>Migrated 50M+ records with 0 data loss by leading design reviews for 7 AML datasets to Palantir OSv2.</li><li>Cut policy lookup time 60% by building a full-stack internal platform with Java, JavaScript, and REST APIs</li><li>Served 100+ compliance stakeholders with real-time endpoints via Java Spring Boot services on Jenkins pipelines</li><li>Resolved 95% of production defects by diagnosing backend issues and providing hands-on support across the SDLC</li><li>Shipped 30+ features and bug fixes across release cycles via hands-on coding in Java Spring Framework</li><li>Reduced release deployment errors 35% by automating build and test stages for backend components in Jenkins</li></ul> | Jul 2023 – Aug 2024<br>Chennai, IN          |

## PROJECTS

**HemsHappen: RAG Medical Chatbot** | *LangChain · FAISS · QLoRA · SentenceTransformers* | [Link](#)

- Built RAG pipeline with FAISS and Claude Sonnet over curated GI guidelines, validated by gastroenterologists
- Fine-tuned Med42-8B via 4-bit QLoRA on a red-flag/triage task, improving F1 4.2 pts & schema validity to 95.6%

**SplitRight: Expense Splitting PWA** | *Next.js 14 · REST API · PostgreSQL · JWT · PWA · Docker* | [Link](#)

- Built full-stack expense-splitting PWA with Next.js 14, Node.js, and PostgreSQL, deployed via Docker on Vercel
- Achieved 90% test coverage across 3 split modes by implementing unit and integration tests with Jest

## TECHNICAL SKILLS

**Languages:** Python, Java, SQL, JavaScript, TypeScript, C/C++, R  
**Frameworks:** React, Next.js, Node.js, Spring Boot, FastAPI, Flask, LangChain  
**Data & ML:** Apache Spark, Hadoop, Prophet, RAG, FAISS, QLoRA  
**Cloud & DevOps:** AWS (S3, EC2), Palantir Foundry, Docker, Jenkins, GitHub Actions, CI/CD  
**Practices:** SDLC, Agile/Scrum, Code review

## CERTIFICATIONS

|                                                                                          |                     |
|------------------------------------------------------------------------------------------|---------------------|
| <b>AWS Certified AI Practitioner</b>   <i>Amazon Web Services</i>   <a href="#">Link</a> | Jul 2026 – Jul 2029 |
|------------------------------------------------------------------------------------------|---------------------|